

Capstone Project - Coronary Heart Disease Study

Sanju Hyacinth C

29/12/2019

```
# PACKAGES REQUIRED:
# install.packages("Hmisc")

## Data loading

setwd("D:/R Progms/CAPSTONE")
getwd()

## [1] "D:/R Progms/CAPSTONE"

heart.d = read.csv("Coronary_heart_risk_study.csv")
#View(heart.d)

head(heart.d)

##   male age education currentSmoker cigsPerDay BPMeds prevalentStroke
## 1    1  39         4             0          0      0              0
## 2    0  46         2             0          0      0              0
## 3    1  48         1             1         20      0              0
## 4    0  61         3             1         30      0              0
## 5    0  46         3             1         23      0              0
## 6    0  43         2             0          0      0              0
##   prevalentHyp diabetes totChol sysBP diaBP BMI heartRate glucose
## 1             0        0    195 106.0   70 26.97      80      77
## 2             0        0    250 121.0   81 28.73      95      76
## 3             0        0    245 127.5   80 25.34      75      70
## 4             1        0    225 150.0   95 28.58      65     103
## 5             0        0    285 130.0   84 23.10      85      85
## 6             1        0    228 180.0  110 30.30      77      99
##   TenYearCHD
## 1           0
## 2           0
## 3           0
## 4           1
## 5           0
## 6           0

## Data Structure and Summary

str(heart.d)

## 'data.frame':   4240 obs. of  16 variables:
##  $ male      : int  1 0 1 0 0 0 0 0 1 1 ...
##  $ age       : int  39 46 48 61 46 43 63 45 52 43 ...
##  $ education : int  4 2 1 3 3 2 1 2 1 1 ...
##  $ currentSmoker : int  0 0 1 1 1 0 0 1 0 1 ...
##  $ cigsPerDay  : int  0 0 20 30 23 0 0 20 0 30 ...
##  $ BPMeds     : int  0 0 0 0 0 0 0 0 0 0 ...
##  $ prevalentStroke: int  0 0 0 0 0 0 0 0 0 0 ...
```

```
## $ prevalentHyp : int 0 0 0 1 0 1 0 0 1 1 ...
## $ diabetes : int 0 0 0 0 0 0 0 0 0 0 ...
## $ totChol : int 195 250 245 225 285 228 205 313 260 225 ...
## $ sysBP : num 106 121 128 150 130 ...
## $ diaBP : num 70 81 80 95 84 110 71 71 89 107 ...
## $ BMI : num 27 28.7 25.3 28.6 23.1 ...
## $ heartRate : int 80 95 75 65 85 77 60 79 76 93 ...
## $ glucose : int 77 76 70 103 85 99 85 78 79 88 ...
## $ TenYearCHD : int 0 0 0 1 0 0 1 0 0 0 ...
```

```
# Our target variable is binary, hence to be converted to factor
# Our continuous numeric terms to be converted from chr to num
# And the nominal terms from num to factor
```

```
## Structure conversion:
```

```
heart.d$male = as.factor(heart.d$male)
heart.d$education = as.factor(heart.d$education)
heart.d$currentSmoker = as.factor(heart.d$currentSmoker)
heart.d$BPMeds = as.factor(heart.d$BPMeds)
heart.d$prevalentStroke = as.factor(heart.d$prevalentStroke)
heart.d$prevalentHyp = as.factor(heart.d$prevalentHyp)
heart.d$diabetes = as.factor(heart.d$diabetes)
heart.d$TenYearCHD = as.factor(heart.d$TenYearCHD)
```

```
## Summary:
```

```
summary(heart.d)
```

```
## male age education currentSmoker cigsPerDay
## 0:2420 Min. :32.00 1 :1720 0:2145 Min. : 0.000
## 1:1820 1st Qu.:42.00 2 :1253 1:2095 1st Qu.: 0.000
## Median :49.00 3 : 689 Median : 0.000
## Mean :49.58 4 : 473 Mean : 9.006
## 3rd Qu.:56.00 NA's: 105 3rd Qu.:20.000
## Max. :70.00 Max. :70.000
## NA's :29
## BPMeds prevalentStroke prevalentHyp diabetes totChol
## 0 :4063 0:4215 0:2923 0:4131 Min. :107.0
## 1 : 124 1: 25 1:1317 1: 109 1st Qu.:206.0
## NA's: 53 Median :234.0
## Mean :236.7
## 3rd Qu.:263.0
## Max. :696.0
## NA's :50
## sysBP diaBP BMI heartRate
## Min. : 83.5 Min. : 48.0 Min. :15.54 Min. : 44.00
## 1st Qu.:117.0 1st Qu.: 75.0 1st Qu.:23.07 1st Qu.: 68.00
## Median :128.0 Median : 82.0 Median :25.40 Median : 75.00
## Mean :132.4 Mean : 82.9 Mean :25.80 Mean : 75.88
## 3rd Qu.:144.0 3rd Qu.: 90.0 3rd Qu.:28.04 3rd Qu.: 83.00
## Max. :295.0 Max. :142.5 Max. :56.80 Max. :143.00
## NA's :19 NA's :1
## glucose TenYearCHD
## Min. : 40.00 0:3596
```

```
## 1st Qu.: 71.00    1: 644
## Median : 78.00
## Mean   : 81.96
## 3rd Qu.: 87.00
## Max.   :394.00
## NA's   :388
```

```
# From observing the target variable, we find that close to 85% of the patients have been cleared of ha
# More than 50% of the population is women, non current smokers and not under blood pressure medication
```

```
## NA counts:
```

```
anyNA(heart.d)
```

```
## [1] TRUE
```

```
# A total of 7 fields have NA values with heartrate having the least and glucose having the most
# We will have to impute data for glucose values as it is very important
# We can remove the NA values in heartrate and education as it will not be a significant data loss
#
```

```
# UNIVARIATE ANALYSIS:
```


```
# Age:
```

```
summary(heart.d$age)
```

```
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
##    32.00  42.00   49.00   49.58  56.00   70.00
```

```
boxplot(heart.d$age, data = heart.d, col = "lightblue", main = "Age distribution")
```

Age distribution



```
# No outliers. Well within the ranges
```


```
# Cigarettes per day
```

```
summary(heart.d$cigsPerDay)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's  
##    0.000   0.000   0.000   9.006  20.000  70.000     29
```

```
boxplot(heart.d$cigsPerDay, data = heart.d, col = "pink", main = "cigsPerDay distribution")
```

cigsPerDay distribution



```
# 2 outliers present
```


```
# Total Cholesterol:
```

```
summary(heart.d$totChol)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's  
##    107.0   206.0   234.0   236.7   263.0   696.0     50
```

```
boxplot(heart.d$totChol, data = heart.d, col = "aquamarine3", main = "Cholesterol distribution")
```

Cholesterol distribution




```
# Systolic BP
summary(heart.d$sysBP)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      83.5   117.0   128.0   132.4   144.0   295.0

boxplot(heart.d$sysBP, data = heart.d, col = "lemonchiffon", main = "Systolic BP distribution")
```

Systolic BP distribution




```
# Diastolic BP
```

```
summary(heart.d$diaBP)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      48.0   75.0   82.0   82.9   90.0   142.5
```

```
boxplot(heart.d$diaBP, data = heart.d, col = "violet", main = "Diastolic BP distribution")
```

Diastolic BP distribution




```
# Body Mass Index:  
summary(heart.d$BMI)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's  
##    15.54   23.07   25.40   25.80   28.04   56.80        19
```

```
boxplot(heart.d$BMI, data = heart.d, col = "orange", main = "BMI distribution")
```


BMI distribution




```
# Heart Rate:
```

```
summary(heart.d$heartRate)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's  
##    44.00   68.00   75.00   75.88   83.00   143.00         1
```

```
boxplot(heart.d$heartRate, data = heart.d, col = "lightgreen", main = "Heart Rate distribution")
```

Heart Rate distribution



```
# Glucose
summary(heart.d$glucose)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's
##    40.00  71.00   78.00   81.96  87.00   394.00     388

boxplot(heart.d$glucose, data = heart.d, col = "tomato", main = "Glucose distribution")

table(heart.d$male, heart.d$TenYearCHD)

##
##          0      1
##    0 2119   301
##    1 1477   343

table(heart.d$male, heart.d$currentSmoker)

##
##          0      1
##    0 1431   989
##    1   714 1106

# Women: 2420, Men: 1820
# Most women (1431 - 59%) are non-smokers while the others (989) are smokers.
# Most of the men (1106 - 60%) are smokers while the others (714) are non-smokers.

table(heart.d$male, heart.d$BPMeds)

##
```

```
##           0      1
##    0 2293    89
##    1 1770    35
```

Close to 95% of the women and 98% of the men do not take BP medication

```
table(heart.d$male, heart.d$prevalentStroke)
```

```
##
##           0      1
##    0 2405    15
##    1 1810    10
```

Almost the entire patient population is free of having had prevalent Stroke

```
table(heart.d$male, heart.d$prevalentHyp)
```


```
##
##           0      1
##    0 1674   746
##    1 1249   571
```

Almost 69% of the female popultion and male population have not suffered prevalent Hypertension

Histograms on ggplot2:

```
library(ggplot2)
```

Glucose distribution




```

agep1 = ggplot(data = heart.d) + geom_histogram(mapping = aes(x = heart.d$age, fill = heart.d$TenYearCHD
cigsp2 = ggplot(data = heart.d) + geom_histogram(mapping = aes(x = heart.d$cigsPerDay, fill = heart.d$T
syp3 = ggplot(data = heart.d) + geom_histogram(mapping = aes(x = heart.d$sysBP, fill = heart.d$TenYear
diap4 = ggplot(data = heart.d) + geom_histogram(mapping = aes(x = heart.d$diaBP, fill = heart.d$TenYear
BMIp5 = ggplot(data = heart.d) + geom_histogram(mapping = aes(x = heart.d$BMI, fill = heart.d$TenYearCHD
heartRp6 = ggplot(data = heart.d) + geom_histogram(mapping = aes(x = heart.d$heartRate, fill = heart.d$
glucp7 = ggplot(data = heart.d) + geom_histogram(mapping = aes(x = heart.d$glucose, fill = heart.d$TenY
totchop8 = ggplot(data = heart.d) + geom_histogram(mapping = aes(x = heart.d$totChol, fill = heart.d$Ten

# View charts
agep1

```

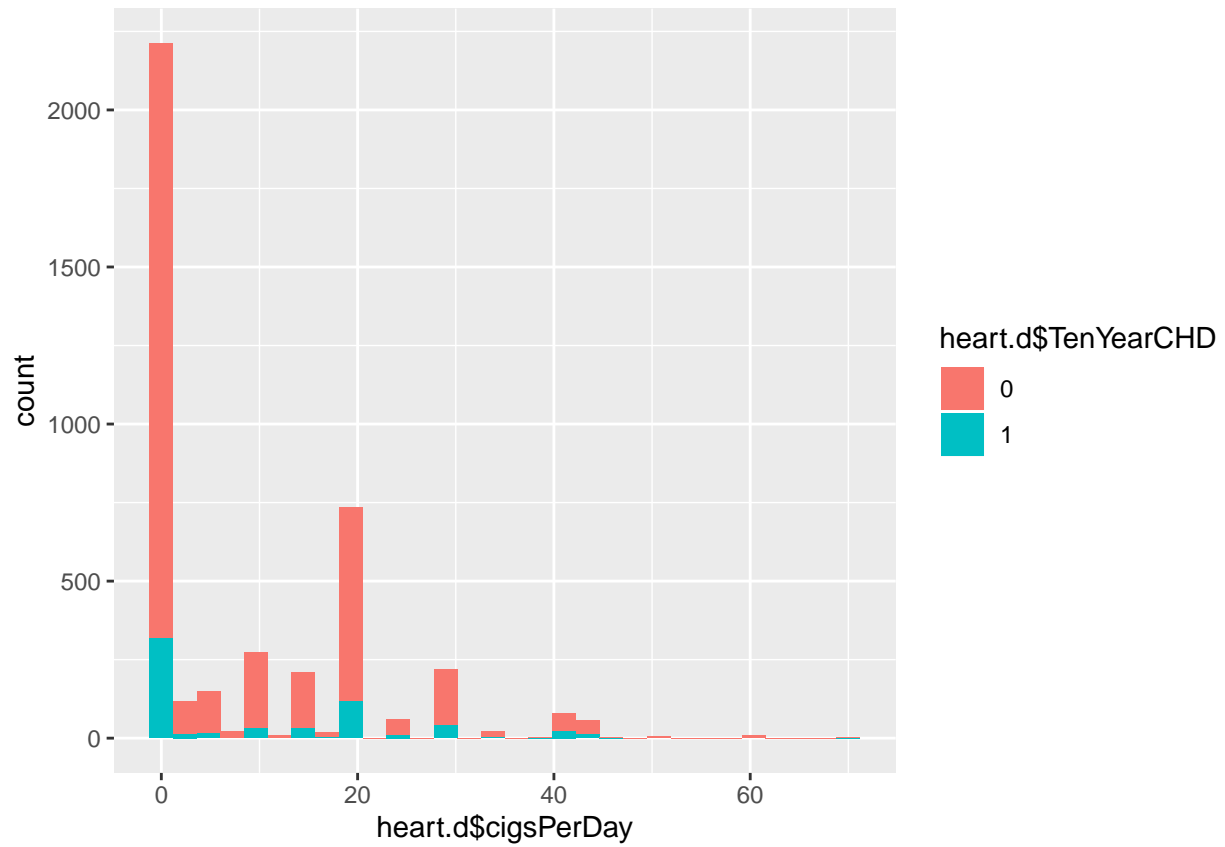
```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
cigsp2
```

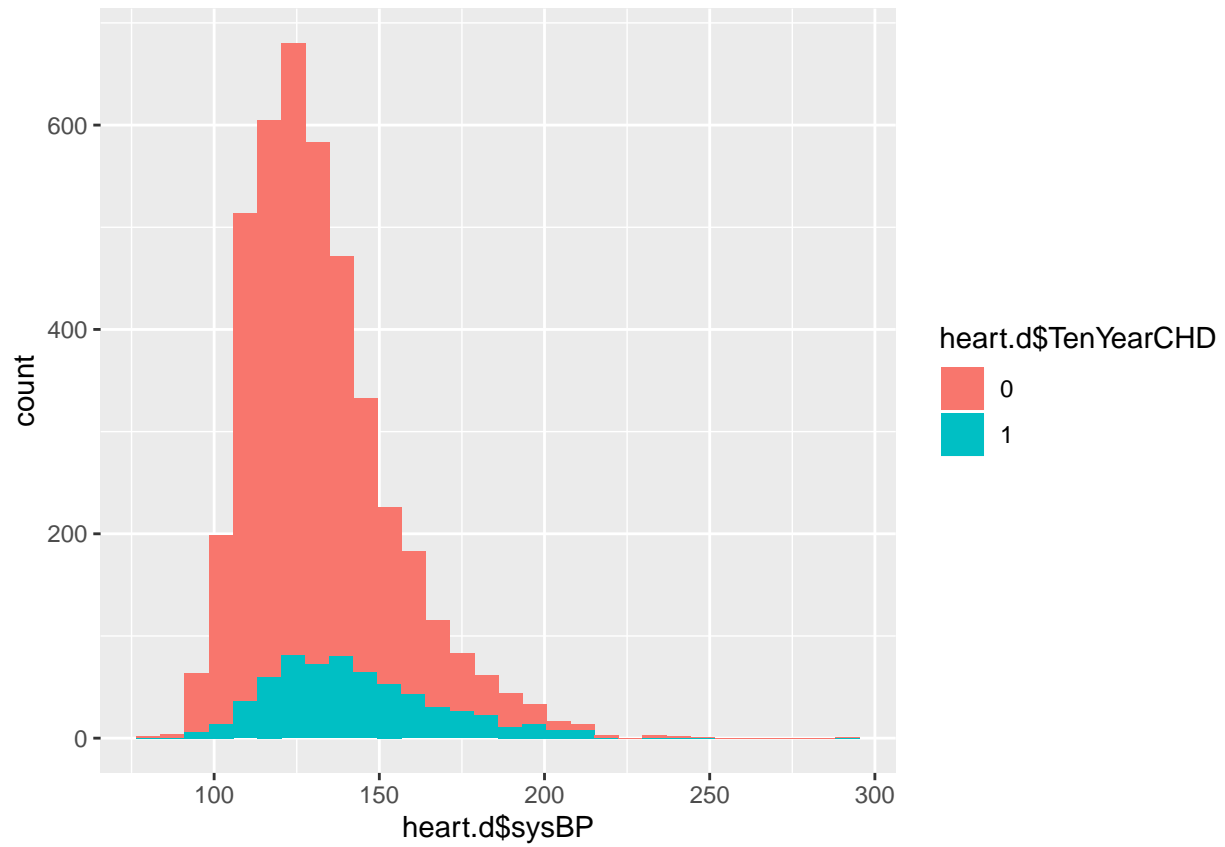
```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

```
## Warning: Removed 29 rows containing non-finite values (stat_bin).
```




syp3

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




diap4

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



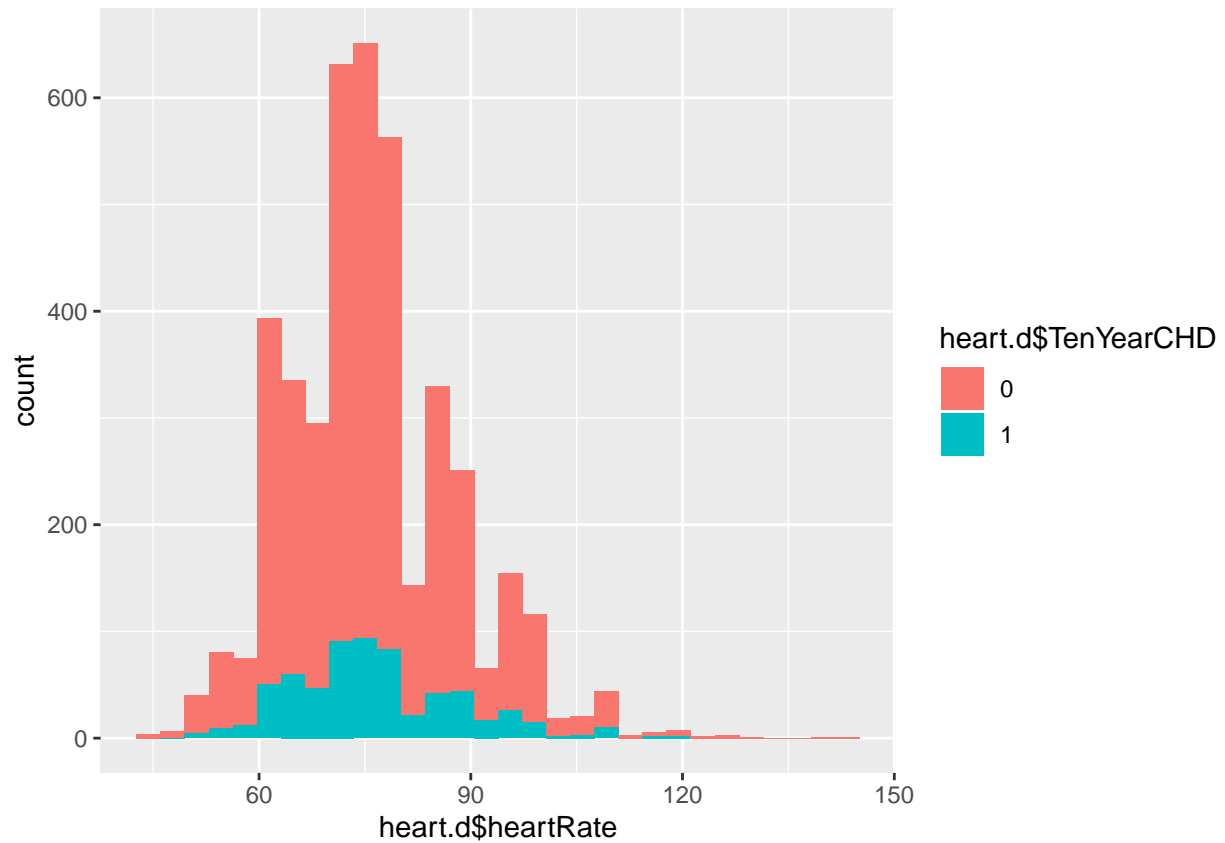
BMIp5

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.  
## Warning: Removed 19 rows containing non-finite values (stat_bin).
```




```
heartRp6
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.  
## Warning: Removed 1 rows containing non-finite values (stat_bin).
```


```
glucp7
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.  
## Warning: Removed 388 rows containing non-finite values (stat_bin).
```



```
totchop8
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.  
## Warning: Removed 50 rows containing non-finite values (stat_bin).
```



BIVARIATE ANALYSIS:

```
library(ggplot2)

# age and cholesterol
agechol1 = ggplot(data = heart.d, aes(x = heart.d$age,
y = heart.d$totChol ,color = heart.d$TenYearCHD)) + geom_point()

# SysBP and DiaBP
sysdia2 = ggplot(data = heart.d, aes(x = heart.d$sysBP,
y = heart.d$diaBP ,color = heart.d$TenYearCHD)) + geom_point()

# Cigs count and age
Cigsage3 = ggplot(data = heart.d, aes(x = heart.d$cigsPerDay,
y = heart.d$age ,color = heart.d$TenYearCHD)) + geom_point()

# Prevalent Hyp and Heart rate:
HypHr4 = ggplot(data = heart.d)+ geom_boxplot(aes(x = heart.d$prevalentHyp,
y = heart.d$heartRate,
fill = heart.d$TenYearCHD))


# Prevalent Hyp and Age
HypHr4.2 = ggplot(data = heart.d)+ geom_boxplot(aes(x = heart.d$prevalentHyp,
y = heart.d$age,
fill = heart.d$TenYearCHD))
```

```
# BP Meds and BMI
```


```
BPMBMI5 = ggplot(data = heart.d) + geom_boxplot(aes(x = heart.d$BPMeds,  
                                                    y = heart.d$BMI,  
                                                    fill = heart.d$TenYearCHD))
```

```
agechol1
```

```
## Warning: Removed 50 rows containing missing values (geom_point).
```




```
sysdia2
```




Cigsage3

```
## Warning: Removed 29 rows containing missing values (geom_point).
```




HypHr4

```
## Warning: Removed 1 rows containing non-finite values (stat_boxplot).
```

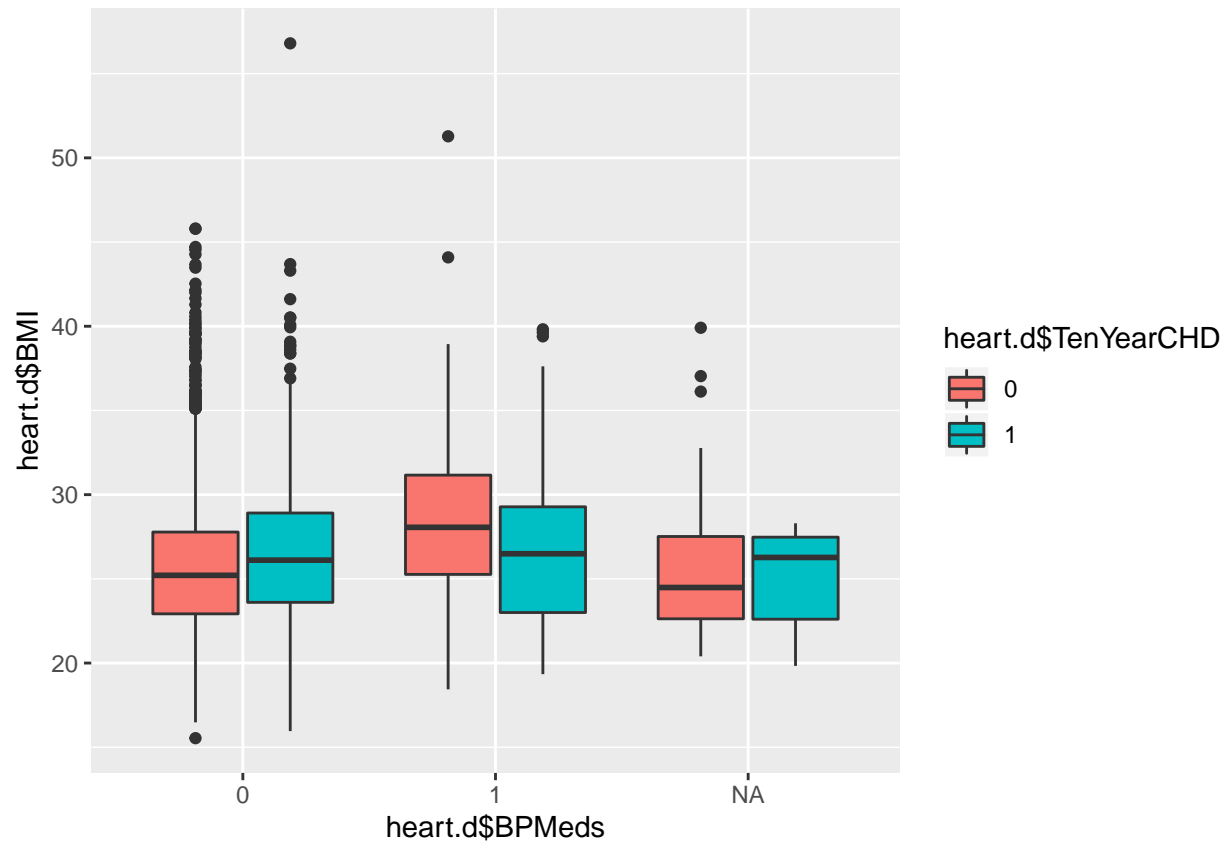


HypHr4.2



BPMBMI5

```
## Warning: Removed 19 rows containing non-finite values (stat_boxplot).
```

```
# Missing value treatment
```

```
# MICE:
library(mice)
```

```
## Warning: package 'mice' was built under R version 3.6.1
```

```
## Loading required package: lattice
```

```
##
```

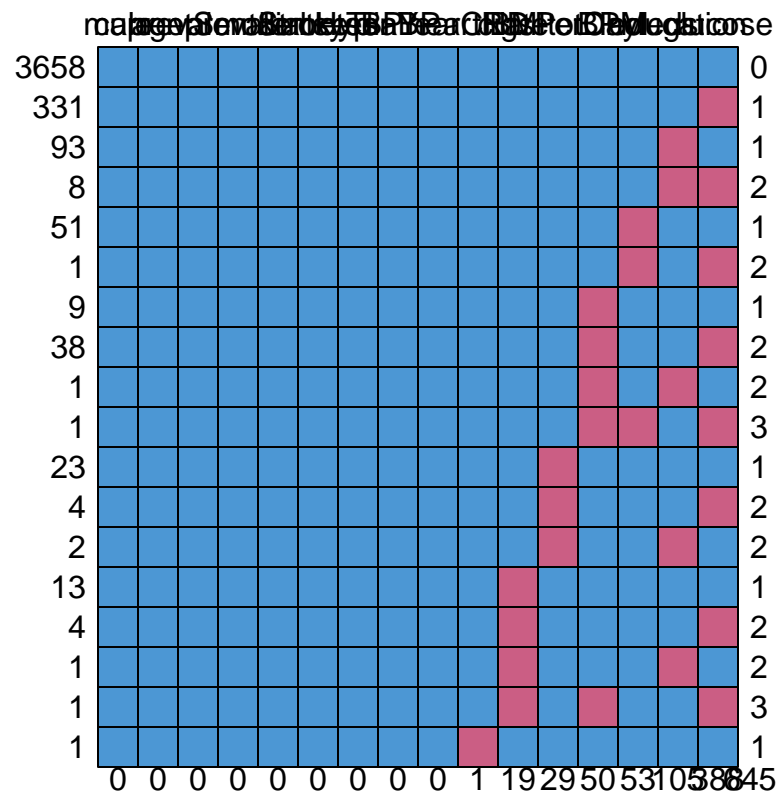
```
## Attaching package: 'mice'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## cbind, rbind
```

```
md.pattern(heart.d)
```



```
##      male age currentSmoker prevalentStroke prevalentHyp diabetes sysBP
## 3658    1   1             1                1              1        1    1
## 331     1   1             1                1              1        1    1
## 93      1   1             1                1              1        1    1
## 8       1   1             1                1              1        1    1
## 51      1   1             1                1              1        1    1
## 1       1   1             1                1              1        1    1
## 9       1   1             1                1              1        1    1
## 38      1   1             1                1              1        1    1
## 1       1   1             1                1              1        1    1
## 1       1   1             1                1              1        1    1
## 23      1   1             1                1              1        1    1
## 4       1   1             1                1              1        1    1
## 2       1   1             1                1              1        1    1
## 13      1   1             1                1              1        1    1
## 4       1   1             1                1              1        1    1
## 1       1   1             1                1              1        1    1
## 1       1   1             1                1              1        1    1
## 1       1   1             1                1              1        1    1
##      0   0             0                0              0        0    0
##      diaBP TenYearCHD heartRate BMI  cigsPerDay totChol BPMeds education
## 3658     1           1         1   1           1       1       1         1
## 331      1           1         1   1           1       1       1         1
## 93       1           1         1   1           1       1       1         0
## 8        1           1         1   1           1       1       1         0
## 51       1           1         1   1           1       1       0         1
```

```
## 1      1      1      1  1      1      1      0      1
## 9      1      1      1  1      1      0      1      1
## 38     1      1      1  1      1      0      1      1
## 1      1      1      1  1      1      0      1      0
## 1      1      1      1  1      1      0      0      1
## 23     1      1      1  1      0      1      1      1
## 4      1      1      1  1      0      1      1      1
## 2      1      1      1  1      0      1      1      0
## 13     1      1      1  1      1      1      1      1
## 4      1      1      1  1      1      1      1      1
## 1      1      1      1  1      1      1      1      0
## 1      1      1      1  1      0      1      1      1
## 1      1      1      0  1      1      1      1      1
##        0      0      1  19     29     50     53     105
##      glucose
## 3658     1  0
## 331      0  1
## 93       1  1
## 8        0  2
## 51       1  1
## 1        0  2
## 9        1  1
## 38       0  2
## 1        1  2
## 1        0  3
## 23       1  1
## 4        0  2
## 2        1  2
## 13       1  1
## 4        0  2
## 1        1  2
## 1        0  3
## 1        1  1
##      388 645
```

```
# 1. REMOVING NA VALUES:
dataforNA = heart.d
dataforNA = na.omit(dataforNA)
1-(3658/4240)
```

```
## [1] 0.1372642
```

```
# 86.27% data is only clean
# 13.73% data is missing
```

```
anyNA(dataforNA)
```

```
## [1] FALSE
```

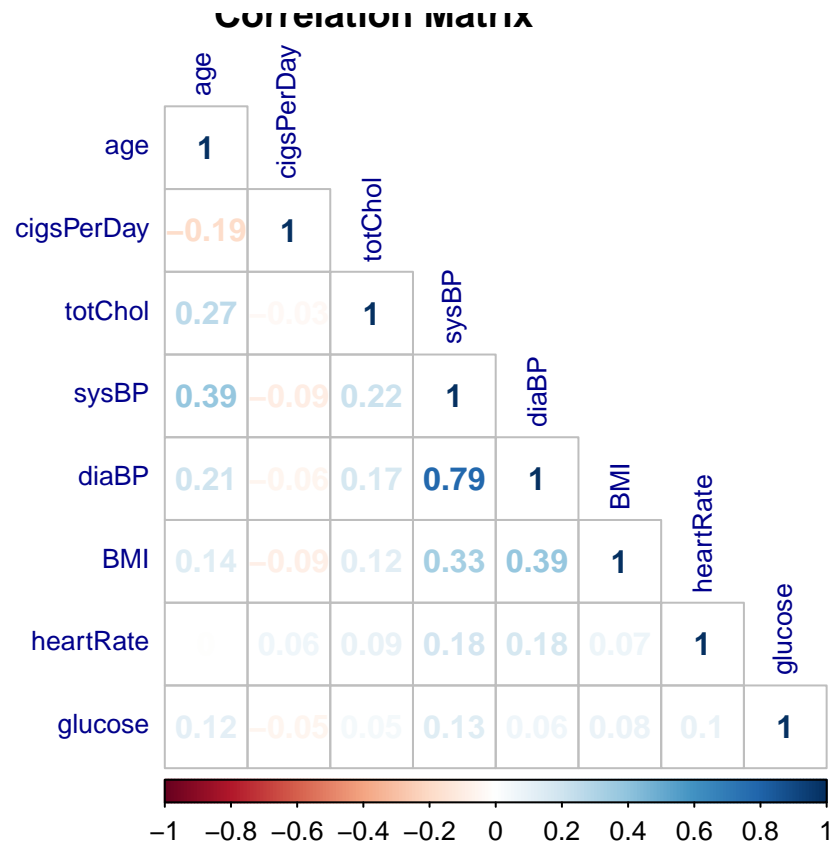
CORRELATION:

```
library(corrplot)
```

```
## Warning: package 'corrplot' was built under R version 3.6.2
```

```
## corrplot 0.84 loaded
```

```
heartcorr = corrplot(cor(dataforNA[, c(2,5, 10:15)]), method = "number", type = "lower", title = "Correlation Matrix")
```



```
#Pairwise correlation
```

```
library(ppcor)
```

```
## Warning: package 'ppcor' was built under R version 3.6.1
```

```
## Loading required package: MASS
```

```
pcor(dataforNA[, c(2,5, 10:15)], method = "pearson")
```

```
## $estimate
```

```
##          age  cigsPerDay  totChol  sysBP  diaBP
## age      1.0000000 -0.154835298 0.210205060 0.34076598 -0.166223048
## cigsPerDay -0.15483530 1.000000000 0.022591777 -0.01798002 0.008122947
## totChol    0.21020506 0.022591777 1.000000000 0.04379225 0.024858780
## sysBP      0.34076598 -0.017980017 0.043792252 1.00000000 0.749437970
## diaBP     -0.16622305 0.008122947 0.024858780 0.74943797 1.000000000
## BMI        0.02543646 -0.059999728 0.044952785 0.01798060 0.220981347
## heartRate  -0.07965411 0.072165033 0.069386258 0.08378496 0.042539111
## glucose    0.05768621 -0.032898880 0.001038711 0.09550652 -0.071830119
##          BMI  heartRate  glucose
## age      0.025436465 -0.079654107 0.057686208
## cigsPerDay -0.059999728 0.072165033 -0.032898880
## totChol    0.044952785 0.069386258 0.001038711
## sysBP      0.017980596 0.083784957 0.095506519
```

```

## diaBP      0.220981347  0.042539111 -0.071830119
## BMI        1.000000000  0.002510351  0.052899322
## heartRate  0.002510351  1.000000000  0.084876926
## glucose    0.052899322  0.084876926  1.000000000
##
## $p.value
##           age  cigsPerDay  totChol  sysBP
## age      0.000000e+00 4.934421e-21 9.481537e-38 5.673968e-100
## cigsPerDay 4.934421e-21 0.000000e+00 1.722631e-01 2.773550e-01
## totChol    9.481537e-38 1.722631e-01 0.000000e+00 8.125559e-03
## sysBP      5.673968e-100 2.773550e-01 8.125559e-03 0.000000e+00
## diaBP      4.880701e-24 6.236215e-01 1.331028e-01 0.000000e+00
## BMI        1.243195e-01 2.857163e-04 6.587258e-03 2.773396e-01
## heartRate  1.438097e-06 1.269713e-05 2.707628e-05 3.968599e-07
## glucose    4.870281e-04 4.681220e-02 9.499658e-01 7.343429e-09
##           diaBP      BMI  heartRate  glucose
## age      4.880701e-24 1.243195e-01 1.438097e-06 4.870281e-04
## cigsPerDay 6.236215e-01 2.857163e-04 1.269713e-05 4.681220e-02
## totChol    1.331028e-01 6.587258e-03 2.707628e-05 9.499658e-01
## sysBP      0.000000e+00 2.773396e-01 3.968599e-07 7.343429e-09
## diaBP      0.000000e+00 1.238622e-41 1.014059e-02 1.393080e-05
## BMI        1.238622e-41 0.000000e+00 8.794604e-01 1.384107e-03
## heartRate  1.014059e-02 8.794604e-01 0.000000e+00 2.794874e-07
## glucose    1.393080e-05 1.384107e-03 2.794874e-07 0.000000e+00
##
## $statistic
##           age cigsPerDay  totChol  sysBP  diaBP
## age      0.000000 -9.4685986 12.98981429 21.898097 -10.1840828
## cigsPerDay -9.468599 0.0000000 1.36523582 -1.086442 0.4907659
## totChol    12.989814 1.3652358 0.00000000 2.648260 1.5023131
## sysBP      21.898097 -1.0864425 2.64825958 0.000000 68.3872629
## diaBP     -10.184083 0.4907659 1.50231314 68.387263 0.0000000
## BMI        1.537247 -3.6314398 2.71858099 1.086477 13.6890593
## heartRate  -4.827661 4.3712642 4.20211438 5.079749 2.5723386
## glucose    3.490939 -1.9886699 0.06275398 5.796545 -4.3508720
##           BMI  heartRate  glucose
## age      1.5372472 -4.8276607 3.49093878
## cigsPerDay -3.6314398 4.3712642 -1.98866991
## totChol    2.7185810 4.2021144 0.06275398
## sysBP      1.0864775 5.0797486 5.79654545
## diaBP     13.6890593 2.5723386 -4.35087196
## BMI        0.0000000 0.1516639 3.20040573
## heartRate  0.1516639 0.0000000 5.14643016
## glucose    3.2004057 5.1464302 0.00000000
##
## $n
## [1] 3658
##
## $gp
## [1] 6
##
## $method
## [1] "pearson"

```

*# The major correlation (78%) is between Systolic and Diastolic BP
A minor correlation between SysBP and Age, and DiaBP and BMI.*

OUTLIERS TREATMENT:

```
treatOut <- function(x) {
  quant <- quantile(x, probs=c(.25, .75), na.rm = T)
  cap <- quantile(x, probs=c(.05, .95), na.rm = T)
  D <- 1.5 * IQR(x, na.rm = T)
  x[ x < (quant[1] - D) ] <- cap[1]
  x[ x > (quant[2] + D) ] <- cap[2]
  return(x)
}
```

treating the outliers with the function above. It helps to replace the lower 25th percentile with the

```
dataforNA$cigsPerDay= treatOut(dataforNA$cigsPerDay)
summary(dataforNA$cigsPerDay)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    0.000  0.000   0.000   8.941  20.000   50.000
```

```
dataforNA$totChol= treatOut(dataforNA$totChol)
summary(dataforNA$totChol)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##     124    206    234    236    263    347
```

```
dataforNA$sysBP= treatOut(dataforNA$sysBP)
summary(dataforNA$sysBP)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##     83.5   117.0   128.0   131.6   143.9   184.0
```

```
dataforNA$diaBP= treatOut(dataforNA$diaBP)
summary(dataforNA$diaBP)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##     53.00   75.00   82.00   82.64   90.00   112.50
```

```
dataforNA$BMI= treatOut(dataforNA$BMI)
summary(dataforNA$BMI)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##     15.96   23.08   25.38   25.64   28.04   35.42
```

```
dataforNA$heartRate= treatOut(dataforNA$heartRate)
summary(dataforNA$heartRate)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##     47.00   68.00   75.00   75.42   82.00   103.00
```


```
dataforNA$glucose= treatOut(dataforNA$glucose)
summary(dataforNA$glucose)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##     47.00   71.00   78.00   79.74   87.00   111.00
```

```
# BOXPLOTS:
```

```
# Cigarettes per day
```

```
boxplot(dataforNA$cigsPerDay, data = dataforNA, col = "pink", main = "cigsPerDay distribution")
```




```
# 2 outliers present
```

```
# Total Cholesterol:
```


```
boxplot(dataforNA$totChol, data = dataforNA, col = "aquamarine3", main = "Cholesterol distribution")
```

Cholesterol distribution




```
# Systolic BP  
boxplot(dataforNA$sysBP, data = dataforNA, col = "lemonchiffon", main = "Systolic BP distribution")
```


Systolic BP distribution




```
# Diastolic BP  
boxplot(dataforNA$diaBP, data = dataforNA, col = "violet", main = "Diastolic BP distribution")
```

Diastolic BP distribution




```
# Body Mass Index:  
boxplot(dataforNA$BMI, data = dataforNA, col = "orange", main = "BMI distribution")
```

BMI distribution




```
# Heart Rate:  
boxplot(dataforNA$heartRate, data = dataforNA, col = "lightgreen", main = "Heart Rate distribution")
```

Heart Rate distribution




```
# Glucose  
boxplot(dataforNA$glucose, data = dataforNA, col = "tomato", main = "Glucose distribution")
```

Glucose distribution



```
# Education and glucose:
EdGluc6 = ggplot(data = dataforNA) + geom_boxplot(aes(x = dataforNA$education,
                                                       y = dataforNA$totChol,
                                                       fill = dataforNA$TenYearCHD))

EdGluc6
```



*# We should make a change in the approach not to include or manipulate the NA values.
 # Pure data is available (about 88%). So let us go with the available data.
 # We should build a model with the manipulated values as well as the removed values.*

```
library(car)
```

```
## Warning: package 'car' was built under R version 3.6.1
```

```
## Loading required package: carData
```

```
## Registered S3 methods overwritten by 'car':
```

```
##   method                      from
```

```
## influence.merMod              lme4
```

```
## cooks.distance.influence.merMod lme4
```

```
## dfbeta.influence.merMod       lme4
```

```
## dfbetas.influence.merMod      lme4
```

```
library(caTools)
```

```
## Warning: package 'caTools' was built under R version 3.6.1
```

```
set.seed(248)
```

```
samplNA = sample.split(dataforNA, SplitRatio = 0.75)
```

```
wNAtrain = subset(dataforNA, samplNA == TRUE)
```

```
wNAtest = subset(dataforNA, samplNA == FALSE)
```

```
prop.table(table(wNAtest$TenYearCHD))
```

```
##
```

```
##           0           1
## 0.8535519 0.1464481

prop.table(table(wNAtrain$TenYearCHD))

##
##           0           1
## 0.8457893 0.1542107

mod1 = glm(TenYearCHD~. , data = wNAtrain, family = "binomial")
summary(mod1)

##
## Call:
## glm(formula = TenYearCHD ~ ., family = "binomial", data = wNAtrain)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6626  -0.5950  -0.4200  -0.2775   2.9039
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -8.233475   0.919343  -8.956 < 2e-16 ***
## male1         0.485478   0.126960   3.824 0.000131 ***
## age           0.067377   0.007750   8.694 < 2e-16 ***
## education2    -0.159283   0.143141  -1.113 0.265807
## education3    -0.100605   0.170122  -0.591 0.554275
## education4    -0.092057   0.193485  -0.476 0.634228
## currentSmoker1 -0.063816   0.185302  -0.344 0.730556
## cigsPerDay     0.025313   0.007547   3.354 0.000796 ***
## BPMeds1       -0.033053   0.280175  -0.118 0.906090
## prevalentStroke1 1.923565   0.676970   2.841 0.004491 **
## prevalentHyp1  0.178703   0.167020   1.070 0.284640
## diabetes1      0.715954   0.289083   2.477 0.013263 *
## totChol        0.001653   0.001428   1.158 0.246814
## sysBP          0.022811   0.005150   4.429 9.45e-06 ***
## diaBP          -0.014704   0.007994  -1.840 0.065840 .
## BMI            0.004932   0.016870   0.292 0.770001
## heartRate      -0.001299   0.005245  -0.248 0.804407
## glucose        0.003280   0.004630   0.708 0.478642
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2358.7  on 2742  degrees of freedom
## Residual deviance: 2067.5  on 2725  degrees of freedom
## AIC: 2103.5
##
## Number of Fisher Scoring iterations: 5

# Prediction on test data:
# Let us predict on the test with mod1 logistic model

test.predict1 = predict(mod1, newdata = wNAtest, type = "response")
table(wNAtest$TenYearCHD, test.predict1>0.5)
```

```
##
##      FALSE TRUE
##    0    769   12
##    1    127    7

# overall accuracy - 84.81%
# specificity - 98.46%
# 5.2% sensitivity
# 2103.5 -> AIC

#      FALSE TRUE
#    0    769   12
#    1    127    7

mod2 = glm(TenYearCHD~male+age+sysBP+cigsPerDay+prevalentStroke+diabetes, data = wNAtrain, family = "binomial")
summary(mod2)

##
## Call:
## glm(formula = TenYearCHD ~ male + age + sysBP + cigsPerDay +
##      prevalentStroke + diabetes, family = "binomial", data = wNAtrain)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.7014  -0.5972  -0.4227  -0.2811   2.8625
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -8.692898   0.500445 -17.370 < 2e-16 ***
## male1         0.460567   0.121493   3.791 0.00015 ***
## age           0.073153   0.007373   9.922 < 2e-16 ***
## sysBP         0.019959   0.002837   7.035 1.99e-12 ***
## cigsPerDay    0.023060   0.004928   4.679 2.88e-06 ***
## prevalentStroke1 1.895331   0.668940   2.833 0.00461 **
## diabetes1     0.820029   0.268703   3.052 0.00227 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2358.7  on 2742  degrees of freedom
## Residual deviance: 2074.9  on 2736  degrees of freedom
## AIC: 2088.9
##
## Number of Fisher Scoring iterations: 5

test.predict2 = predict(mod2, newdata = wNAtest, type = "response")
table(wNAtest$TenYearCHD, test.predict2>0.5)

##
##      FALSE TRUE
##    0    769   12
##    1    127    7

# Not a change in the model is seen. The values are just the same.
# 2088.9 -> AIC
```



```

# doing smote

library(DMwR)

## Warning: package 'DMwR' was built under R version 3.6.1
## Loading required package: grid
## Registered S3 method overwritten by 'xts':
##   method      from
##   as.zoo.xts zoo

## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo

sm.train = subset(dataforNA, samplNA == TRUE)
sm.test = subset(dataforNA, samplNA == FALSE)

# on the NA set

prop.table(table(sm.train$TenYearCHD))

##
##           0           1
## 0.8457893 0.1542107

prop.table(table(sm.test$TenYearCHD))

##
##           0           1
## 0.8535519 0.1464481

balanced.train = SMOTE(TenYearCHD~., sm.train, perc.over = 100, k = 5, perc.under = 400)
table(balanced.train$TenYearCHD)


##
##      0      1
## 1692  846

prop.table(table(balanced.train$TenYearCHD))

##
##           0           1
## 0.6666667 0.3333333

barplot(prop.table(table(balanced.train$TenYearCHD)))

```



we have like 66 to 33 percentage now

```
library(car)
```

```
mod3 = glm(TenYearCHD~. , data = balanced.train, family = "binomial")
summary(mod3)
```

```
##
## Call:
## glm(formula = TenYearCHD ~ ., family = "binomial", data = balanced.train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4681  -0.7826  -0.5168   0.8738   2.7006
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -8.868449   0.774217 -11.455  < 2e-16 ***
## male1         0.383538   0.105083   3.650 0.000262 ***
## age          0.061691   0.006754   9.134  < 2e-16 ***
## education2   -0.209279   0.120403  -1.738 0.082184 .
## education3   -0.156604   0.140596  -1.114 0.265340
## education4   -0.092581   0.159704  -0.580 0.562114
## currentSmoker1 0.341686   0.134928   2.532 0.011330 *
## cigsPerDay    0.009243   0.005916   1.562 0.118187
## BPMeds1       1.271206   0.204256   6.224 4.86e-10 ***
## prevalentStroke1 1.868723   0.439452   4.252 2.11e-05 ***
```

```
## prevalentHyp1      0.002853    0.130039    0.022 0.982499
## diabetes1          2.003430    0.256585    7.808 5.81e-15 ***
## totChol            0.002451    0.001241    1.975 0.048243 *
## sysBP              0.022239    0.004413    5.040 4.66e-07 ***
## diaBP             -0.008034    0.007143   -1.125 0.260705
## BMI                0.035924    0.014749    2.436 0.014862 *
## heartRate         -0.002859    0.004522   -0.632 0.527210
## glucose            0.009055    0.003940    2.298 0.021551 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 3230.9  on 2537  degrees of freedom
## Residual deviance: 2608.8  on 2520  degrees of freedom
## AIC: 2644.8
##
## Number of Fisher Scoring iterations: 4
```

```
vif(mod3)
```

```
##              GVIF Df GVIF^(1/(2*Df))
## male          1.190536  1      1.091117
## age           1.271062  1      1.127414
## education     1.100553  3      1.016097
## currentSmoker 1.964845  1      1.401729
## cigsPerDay    2.086939  1      1.444624
## BPMeds        1.048776  1      1.024097
## prevalentStroke 1.011705  1      1.005836
## prevalentHyp  1.793576  1      1.339245
## diabetes      1.068260  1      1.033567
## totChol       1.075378  1      1.037005
## sysBP         3.218069  1      1.793898
## diaBP         2.682250  1      1.637758
## BMI           1.148236  1      1.071558
## heartRate     1.095848  1      1.046827
## glucose       1.099035  1      1.048349
```

```
# Prediction on test data:
```

```
# Let us predict on the test with mod1 logistic model
```

```
test.predict3 = predict(mod3, newdata = sm.test, type = "response")
table(sm.test$TenYearCHD, test.predict3>0.3)
```

```
##
##      FALSE TRUE
##  0    552  229
##  1     49   85
```

```
# MOST SIGNIFICANT *** : age, male, BP meds, prevStroke, diabetes, sysBP
# SIGNIFICANT * : current smoker, totchol, BMI, glucose
```

```
mod4 = glm(TenYearCHD~age+male+BPMeds+prevalentStroke+sysBP+diabetes+currentSmoker, data = balanced.train, family = "binomial")
summary(mod4)
```

```
##
```

```
## Call:
## glm(formula = TenYearCHD ~ age + male + BPMeds + prevalentStroke +
##      sysBP + diabetes + currentSmoker, family = "binomial", data = balanced.train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.5421  -0.7783  -0.5338   0.8866   2.5245
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -7.604870    0.436384 -17.427 < 2e-16 ***
## age              0.065048    0.006363  10.223 < 2e-16 ***
## male1           0.435761    0.098597   4.420 9.89e-06 ***
## BPMeds1         1.275555    0.203113   6.280 3.38e-10 ***
## prevalentStroke1 1.859899    0.437129   4.255 2.09e-05 ***
## sysBP           0.021047    0.002620   8.032 9.61e-16 ***
## diabetes1       2.206345    0.248383   8.883 < 2e-16 ***
## currentSmoker1  0.426354    0.100645   4.236 2.27e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 3230.9  on 2537  degrees of freedom
## Residual deviance: 2631.3  on 2530  degrees of freedom
## AIC: 2647.3
##
## Number of Fisher Scoring iterations: 4

test.predict4 = predict(mod4, newdata = sm.test, type = "response")
table(sm.test$TenYearCHD, test.predict4>0.3)

##
##      FALSE TRUE
##  0    563   218
##  1     51    83

#      FALSE TRUE
#  0    563   218
#  1     51    83

# Spec -> 72.09
# Sens -> 61.94
# Over -> 70.60

test.predict4.1 = predict(mod4, newdata = sm.test, type = "response")
table(sm.test$TenYearCHD, test.predict4.1>0.4)

##
##      FALSE TRUE
##  0    652   129
##  1     77    57

## NAIVE BAYES:
library(e1071)
```

```
## Warning: package 'e1071' was built under R version 3.6.1
```

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 3.6.1
```

```
# doing on the omitted NA set:
```

```
NB1 = naiveBayes(TenYearCHD~., data = balanced.train)
```

```
print(NB1)
```

```
##
```

```
## Naive Bayes Classifier for Discrete Predictors
```

```
##
```

```
## Call:
```

```
## naiveBayes.default(x = X, y = Y, laplace = laplace)
```

```
##
```

```
## A-priori probabilities:
```

```
## Y
```

```
##           0           1
```

```
## 0.6666667 0.3333333
```

```
##
```

```
## Conditional probabilities:
```

```
##   male
```

```
## Y           0           1
```

```
## 0 0.5626478 0.4373522
```

```
## 1 0.4657210 0.5342790
```

```
##
```

```
##   age
```

```
## Y      [,1]      [,2]
```

```
## 0 48.91726 8.532417
```

```
## 1 54.72766 7.214810
```

```
##
```

```
##   education
```

```
## Y           1           2           3           4
```

```
## 0 0.3788416 0.3102837 0.1891253 0.1217494
```

```
## 1 0.4917258 0.2434988 0.1560284 0.1087470
```

```
##
```

```
##   currentSmoker
```

```
## Y           0           1
```

```
## 0 0.5254137 0.4745863
```

```
## 1 0.4988180 0.5011820
```

```
##
```

```
##   cigsPerDay
```

```
## Y      [,1]      [,2]
```

```
## 0 8.765366 11.67928
```

```
## 1 9.382491 10.80129
```

```
##
```

```
##   BPMeds
```

```
## Y           0           1
```

```
## 0 0.9751773 0.0248227
```

```
## 1 0.8557920 0.1442080
```

```
##
```

```
##   prevalentStroke
```

```
## Y           0           1
```

```
## 0 0.995862884 0.004137116
```

```
## 1 0.959810875 0.040189125
##
## prevalentHyp
## Y 0 1
## 0 0.7056738 0.2943262
## 1 0.4598109 0.5401891
##
## diabetes
## Y 0 1
## 0 0.98699764 0.01300236
## 1 0.86288416 0.13711584
##
## totChol
## Y [,1] [,2]
## 0 234.3960 41.28831
## 1 245.8707 38.37990
##
## sysBP
## Y [,1] [,2]
## 0 130.2503 18.66421
## 1 143.6089 20.11318
##
## diaBP
## Y [,1] [,2]
## 0 82.30230 10.72804
## 1 87.16569 11.45281
##
## BMI
## Y [,1] [,2]
## 0 25.49467 3.595087
## 1 26.54812 3.460878
##
## heartRate
## Y [,1] [,2]
## 0 75.64190 11.55883
## 1 76.51296 10.37642
##
## glucose
## Y [,1] [,2]
## 0 78.97991 12.26938
## 1 83.17445 13.37521
```

```
NB.pred1 = predict(NB1, sm.test, type = "class" )
table(NB.pred1, sm.test$TenYearCHD, dnn = c("Prediction", "Actual"))
```

```
##           Actual
## Prediction  0   1
##           0 619  70
##           1 162  64
```

```
# overall accuracy - 74.64%
# specificity - 89.84%
# sensitivity - 28.31%
```

```
# this model where we have omitted for has given an approx 40% accuracy of finding risk bearers.
```

```

# The prediction of non risk bearers is good with the dataset that has NA removed (86.10%)

# We have an overall accuracy of 80.84% with only 88% of the data being useful

#           Actual
# Prediction  0    1
#           0  619  70
#           1  162  64

# only on the categorical variables:

NB2 = naiveBayes(TenYearCHD~., data = balanced.train[, -c(2,3,5,10:15)])
print(NB2)

```

```

##
## Naive Bayes Classifier for Discrete Predictors
##
## Call:
## naiveBayes.default(x = X, y = Y, laplace = laplace)
##
## A-priori probabilities:
## Y
##           0           1
## 0.6666667 0.3333333
##
## Conditional probabilities:
##   male
## Y           0           1
## 0 0.5626478 0.4373522
## 1 0.4657210 0.5342790
##
##   currentSmoker
## Y           0           1
## 0 0.5254137 0.4745863
## 1 0.4988180 0.5011820
##
##   BPMeds
## Y           0           1
## 0 0.9751773 0.0248227
## 1 0.8557920 0.1442080
##
##   prevalentStroke
## Y           0           1
## 0 0.995862884 0.004137116
## 1 0.959810875 0.040189125
##
##   prevalentHyp
## Y           0           1
## 0 0.7056738 0.2943262
## 1 0.4598109 0.5401891
##
##   diabetes
## Y           0           1

```

```
## 0 0.98699764 0.01300236
## 1 0.86288416 0.13711584
```

```
NB.pred2 = predict(NB2, sm.test, type = "class" )
table(NB.pred2, sm.test$TenYearCHD, dnn = c("Prediction", "Actual"))
```

```
##           Actual
## Prediction  0    1
##           0 738 114
##           1  43  20
```

```
#           Actual
# Prediction  0    1
#           0 738 114
#           1  43  20
```

```
# overall accuracy - 82.84%
# specificity - 86.62%
# sensitivity - 31.75%
```

```
# random forest:
```

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.6.1
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
## margin
```

```
RFmtry.val = floor(sqrt(ncol(balanced.train)))
```

```
RFmtry.val
```

```
## [1] 4
```

```
# RF1
```

```
RF.m1 = randomForest(TenYearCHD~., data = balanced.train,
                      ntree = 2000, mtry = RFmtry.val, nodesize = 10, importance = TRUE)
```

```
print(RF.m1)
```

```
##
```

```
## Call:
```

```
## randomForest(formula = TenYearCHD ~ ., data = balanced.train, ntree = 2000, mtry = RFmtry.val,
```

```
## Type of random forest: classification
```

```
## Number of trees: 2000
```

```
## No. of variables tried at each split: 4
```

```
##
```

```
## OOB estimate of error rate: 16.67%
```

```
## Confusion matrix:
```

```
## 0 1 class.error
```

```
## 0 1600 92 0.05437352
```



```
## 1 331 515 0.39125296
```

```
# We see an increased OOB rate, but the class error is lowered (39%).
# We also are able to see a good rise in the prediction of the heart risk bearers.
# Let us try to boost this model for a better result.
```

```
# Confusion matrix:
```


```
#      0    1 class.error
# 0 1602  90  0.05319149
# 1  334 512  0.39479905
```

```
# overall -> 83.29%
```

```
# spec -> 94.68%
```

```
# sens -> 60.52%
```

```
plot(RF.m1)
```



```
importance(RF.m1)
```

##	0	1	MeanDecreaseAccuracy	MeanDecreaseGini
## male	37.64950	3.910741	35.70239	10.637036
## age	108.64038	82.337889	129.92255	131.447952
## education	41.49978	2.192878	37.05837	24.146251
## currentSmoker	67.37669	-64.453255	62.67301	38.236650
## cigsPerDay	102.58427	25.760910	111.51891	100.746873
## BPMeds	69.60891	9.351757	67.19038	25.540830
## prevalentStroke	25.97593	21.503032	31.92736	6.376223


```
## prevalentHyp      55.74786 -29.955942          53.94800      21.701717
## diabetes          74.50619  40.108679          77.18216      38.653019
## totChol           59.31384  17.337663          58.11372      63.411348
## sysBP             87.24799  31.341389         100.47808     111.279116
## diaBP             70.11349  18.736832          78.65527      74.437231
## BMI               65.55799  22.393193          66.47565      71.647512
## heartRate         59.95301  19.683526          59.50893      58.890256
## glucose           65.75196  26.565808          66.08974      65.411160
```

from the importance values, we find that age seems to be the most important parameter to predict the

TUNING:

```
t_RF.m1 = tuneRF(x=balanced.train[,-c(16)],y=balanced.train$TenYearCHD, mtrystart = 15, stepfactor= 1.5,
                 ntree= 1700 , improve = 0.0001, nodesize=10, trace=TRUE, plot=TRUE, doBest=TRUE, impor
```

```
## mtry = 3   OOB error = 16.94%
## Searching left ...
## mtry = 2   OOB error = 17.73%
## -0.04651163 1e-04
## Searching right ...
## mtry = 6   OOB error = 16.82%
## 0.006976744 1e-04
## mtry = 12  OOB error = 16.63%
## 0.0117096 1e-04
## mtry = 15  OOB error = 16.75%
## -0.007109005 1e-04
```



```

# tuned model shows the best value at 12, OOB error rate - 16.59%

# REFINED MODEL: 2

RF.m1 = randomForest(TenYearCHD~.,data = balanced.train, ntree = 1700, mtry = 12, nodesize = 10, importanc
print(RF.m1)

##
## Call:
## randomForest(formula = TenYearCHD ~ ., data = balanced.train,          ntree = 1700, mtry = 12, nodesize
##           Type of random forest: classification
##           Number of trees: 1700
## No. of variables tried at each split: 12
##
##           OOB estimate of  error rate: 16.71%
## Confusion matrix:
##      0   1 class.error
## 0 1595  97  0.05732861
## 1  327 519  0.38652482
# Confusion matrix:
#      0   1 class.error
# 0 1596  96  0.05673759
# 1  324 522  0.38297872

# overall -> 83.45%
# spec -> 94.33%
# sens -> 61.70%

importance(RF.m1)

##
##           0           1 MeanDecreaseAccuracy
## male      28.92889    1.7241395          28.57315
## age      123.53177   75.6453917          148.41907
## education  41.11019    1.8224797          37.58902
## currentSmoker 97.02103 -99.2640007          85.04859
## cigsPerDay 158.76259  10.8188352          160.49151
## BPMeds     61.11769    0.8321436          58.79695
## prevalentStroke 21.75942  21.2515871          30.29471
## prevalentHyp 36.72438 -25.8615770          35.52930
## diabetes    90.62315   31.1861887          89.42834
## totChol     60.62424   12.7192657          59.78540
## sysBP       83.84634   32.8960302          97.42501
## diaBP       70.89521   24.0066777          83.49570
## BMI         70.77156   24.3861981          73.48797
## heartRate   57.86200   14.9250877          56.41450
## glucose     67.91891   23.9712715          69.24409
##           MeanDecreaseGini
## male           8.132826
## age          157.975515
## education      23.487557
## currentSmoker   63.133566
## cigsPerDay     113.485575

```

```

## BPMeds                16.723957
## prevalentStroke        6.078960
## prevalentHyp           9.523759
## diabetes               37.143289
## totChol                68.485520
## sysBP                 109.804005
## diaBP                  81.542630
## BMI                    82.736909
## heartRate              60.592978
## glucose                69.418448

balanced.train$RF.Pred = predict(RF.m1, data = balanced.train, type = "class")
balanced.train$RF.Score = 1 - predict(RF.m1, data = balanced.train, type = "prob")[,2]
sm.test$RF.Pred = predict(RF.m1, newdata = sm.test, type = "class")
sm.test$RF.Score = 1 - predict(RF.m1, newdata = sm.test, type = "prob")[,2]

t_devRF <- with(balanced.train, table(TenYearCHD, RF.Pred))
t_devRF

##           RF.Pred
## TenYearCHD    0    1
##           0 1595   97
##           1  327  519

#           RF.Pred
# TenYearCHD    0    1
#           0 1597   95
#           1  329  517

# RF has made a good prediction - 61.11% CORRECT PREDICTION OF RISK BEARERS
# Spec - 94.38%
# Overall Acc - 83.29%

# MODEL PERFORMANCE MEASURES:
library(InformationValue)

## Warning: package 'InformationValue' was built under R version 3.6.1
##
## Attaching package: 'InformationValue'
## The following objects are masked from 'package:caret':
##
##      confusionMatrix, precision, sensitivity, specificity
library(ROCR)

## Warning: package 'ROCR' was built under R version 3.6.1
## Loading required package: gplots
## Warning: package 'gplots' was built under R version 3.6.1
##
## Attaching package: 'gplots'
## The following object is masked from 'package:stats':
##
##      lowess

```

```
library(ineq)
library(ROSE)
```


```
## Warning: package 'ROSE' was built under R version 3.6.1
```

```
## Loaded ROSE 0.0-3
```


```
# RF:
```

```
#AUC
```

```
auc.train = roc.curve(balanced.train$TenYearCHD, balanced.train$RF.Score)
```



```
auc.test = roc.curve(sm.test$TenYearCHD, sm.test$RF.Score)
```



The AUC is a little over 60% as we can see, and is more than the random prediction.

```
ineq(sm.test$RF.Score,"gini")
```

```
## [1] 0.150951
```

lower gini of 15.01% indicated higher equality or lower inequality in distribution of the risk factors

Boosting

XGBoosting is tried here with only our numeric variables of the NA removed dataset.

```
library(xgboost)
```

```
## Warning: package 'xgboost' was built under R version 3.6.1
```

*# XGBoost works with matrices that contain all numeric variables
we also need to split the training data and label*

```
gd_features_train<-as.matrix(balanced.train[, c(2,5,10:15)])
gd_label_train<-as.matrix(balanced.train[,16])
gd_features_test<-as.matrix(sm.test[,c(2,5,10:15)])
```

```
xgb.fit <- xgboost(
  data = gd_features_train,
  label = gd_label_train,
  eta = 0.001, #this is like shrinkage in the previous algorithm
```

```

max_depth = 3, #Larger the depth, more complex the model; higher chances of overfitting. There is no s
min_child_weight = 3, #it blocks the potential feature interactions to prevent overfitting
nrounds = 1000, #controls the maximum number of iterations. For classification, it is similar to the n
nfold = 5,
objective = "binary:logistic", # for regression models
verbose = 0, # silent,
early_stopping_rounds = 10 # stop if no improvement for 10 consecutive trees
)

#gd_features_test<-as.matrix(gd_features_test[,1:ncol(gd_features_test)-1])

sm.test$xbg.pred.class <- predict(xgb.fit, gd_features_test)

table(sm.test$TenYearCHD, sm.test$xbg.pred.class>0.3)

##
##      TRUE
##  0   781
##  1   134

#this model was definitely better
#or simply the total correct of the minority class

sum(sm.test$TenYearCHD==1 & sm.test$xbg.pred.class>=0.3)

## [1] 134

#      TRUE
#  0   781
#  1   134

#in this code chunk we will playing around with all the values untill we find the best fit
#let's play with shrinkage, known as eta in xgb

tp_xgb<-vector()
lr <- c(0.001, 0.01, 0.1, 0.3, 0.5, 0.7, 1)
md<-c(1,3,5,7,9,15)
nr<-c(2, 50, 100, 1000, 10000)
for (i in md) {

  xgb.fit1 <- xgboost(
    data = gd_features_train,
    label = gd_label_train,
    eta = 0.2,
    max_depth = 15,
    nrounds = 10,
    nfold = 5,
    objective = "binary:logistic", # for regression models
    verbose = 1, # silent,
    early_stopping_rounds = 14 # stop if no improvement for 10 consecutive trees
  )

  sm.test$xbg.pred.class <- predict(xgb.fit1, gd_features_test)

  tp_xgb<-cbind(tp_xgb, sum(sm.test$TenYearCHD==1 & sm.test$xbg.pred.class>=0.3))

```

#if your class=1 and our prediction=0.2, we are going to display it with the next line compare the sa

```
}
```

```
## [1] train-error:0.098897
## Will train until train_error hasn't improved in 14 rounds.
##
## [2] train-error:0.085894
## [3] train-error:0.069740
## [4] train-error:0.057526
## [5] train-error:0.047281
## [6] train-error:0.038219
## [7] train-error:0.034673
## [8] train-error:0.030733
## [9] train-error:0.024823
## [10] train-error:0.020883
## [1] train-error:0.098897
## Will train until train_error hasn't improved in 14 rounds.
##
## [2] train-error:0.085894
## [3] train-error:0.069740
## [4] train-error:0.057526
## [5] train-error:0.047281
## [6] train-error:0.038219
## [7] train-error:0.034673
## [8] train-error:0.030733
## [9] train-error:0.024823
## [10] train-error:0.020883
## [1] train-error:0.098897
## Will train until train_error hasn't improved in 14 rounds.
##
## [2] train-error:0.085894
## [3] train-error:0.069740
## [4] train-error:0.057526
## [5] train-error:0.047281
## [6] train-error:0.038219
## [7] train-error:0.034673
## [8] train-error:0.030733
## [9] train-error:0.024823
## [10] train-error:0.020883
## [1] train-error:0.098897
## Will train until train_error hasn't improved in 14 rounds.
##
## [2] train-error:0.085894
## [3] train-error:0.069740
## [4] train-error:0.057526
## [5] train-error:0.047281
## [6] train-error:0.038219
## [7] train-error:0.034673
## [8] train-error:0.030733
## [9] train-error:0.024823
## [10] train-error:0.020883
## [1] train-error:0.098897
## Will train until train_error hasn't improved in 14 rounds.
```



```
##
## [2] train-error:0.085894
## [3] train-error:0.069740
## [4] train-error:0.057526
## [5] train-error:0.047281
## [6] train-error:0.038219
## [7] train-error:0.034673
## [8] train-error:0.030733
## [9] train-error:0.024823
## [10] train-error:0.020883
## [1] train-error:0.098897
## Will train until train_error hasn't improved in 14 rounds.
##
## [2] train-error:0.085894
## [3] train-error:0.069740
## [4] train-error:0.057526
## [5] train-error:0.047281
## [6] train-error:0.038219
## [7] train-error:0.034673
## [8] train-error:0.030733
## [9] train-error:0.024823
## [10] train-error:0.020883
```

```
tp_xgb
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6]
## [1,]   76   76   76   76   76   76
```

```
table(sm.test$TenYearCHD, sm.test$xgb.pred.class>=0.3)
```

```
##
##      FALSE TRUE
##  0    502  279
##  1     58   76
```

```
# here there is significant improvement of the model compared to our logistic model
# sensitivity is found to be 56.72%, spec 64.27%, overall 63.17% accurate
```

```
# wNAtest = wNAtest[, -17]
```

```
#      FALSE TRUE
#  0    502  279
#  1     58   76
```